

Assignment No-5

Data Analytics II 1. Implement logistic regression using Python/R to perform classification on Social_Network_Ads.csv dataset. 2. Compute Confusion matrix to find TP, FP, TN, FN, Accuracy, Error rate, Precision, Recall on the given dataset

Import libraries and create alias for Pandas, Numpy

```
[1]: import pandas as pd
import numpy as np
```

Import the Social_Media_Adv Dataset

```
[2]: from google.colab import files
```

```
[ ]: files.upload()
```

<IPython.core.display.HTML object>

Saving Social_Network_Ads.csv to Social_Network_Ads (1).csv

```
[ ]: {'Social_Network_Ads (1).csv': b'Age,EstimatedSalary,Purchased\n19,19000,0\n35,20000,0\n26,43000,0\n27,57000,0\n19,76000,0\n27,58000,0\n27,84000,0\n32,150000,1\n25,33000,0\n35,65000,0\n26,80000,0\n26,52000,0\n20,86000,0\n32,18000,0\n18,82000,0\n29,80000,0\n47,25000,1\n45,26000,1\n46,28000,1\n48,29000,1\n45,22000,1\n47,49000,1\n48,41000,1\n45,22000,1\n46,23000,1\n47,20000,1\n49,28000,1\n47,30000,1\n29,43000,0\n31,18000,0\n31,74000,0\n27,137000,1\n21,16000,0\n28,44000,0\n27,90000,0\n35,27000,0\n33,28000,0\n30,49000,0\n26,72000,0\n27,31000,0\n27,17000,0\n33,51000,0\n35,108000,0\n30,15000,0\n28,84000,0\n23,20000,0\n25,79000,0\n27,54000,0\n30,135000,1\n31,89000,0\n24,32000,0\n18,44000,0\n29,83000,0\n35,23000,0\n27,58000,0\n24,55000,0\n23,48000,0\n28,79000,0\n22,18000,0\n32,117000,0\n27,20000,0\n25,87000,0\n23,66000,0\n32,12000,0,1\n59,83000,0\n24,58000,0\n24,19000,0\n23,82000,0\n22,63000,0\n31,68000,0\n25,80000,0\n24,27000,0\n20,23000,0\n33,113000,0\n32,18000,0\n34,112000,1\n18,52000,0\n22,27000,0\n28,87000,0\n26,17000,0\n30,8000,0\n39,42000,0\n20,49000,0\n35,88000,0\n30,62000,0\n31,118000,1\n24,55000,0\n28,85000,0\n26,81000,0\n35,50000,0\n22,81000,0\n30,116000,0
```

```
\n26,15000,0\n29,28000,0\n29,83000,0\n35,44000,0\n35,25000,0\n28,123000,1\n35,73000,0\n28,37000,0\n27,88000,0\n28,59000,0\n32,86000,0\n33,149000,1\n19,21000,0\n21,72000,0\n26,35000,0\n27,89000,0\n26,86000,0\n38,80000,0\n39,71000,0\n37,71000,0\n38,61000,0\n37,55000,0\n42,8000,0\n40,57000,0\n35,75000,0\n36,52000,0\n40,59000,0\n41,59000,0\n36,75000,0\n37,72000,0\n40,75000,0\n35,53000,0\n41,51000,0\n39,61000,0\n42,65000,0\n26,32000,0\n30,17000,0\n26,84000,0\n31,58000,0\n33,31000,0\n30,87000,0\n21,68000,0\n28,55000,0\n23,63000,0\n20,82000,0\n30,107000,1\n28,59000,0\n19,25000,0\n19,85000,0\n18,68000,0\n35,59000,0\n30,89000,0\n34,25000,0\n24,89000,0\n27,96000,1\n41,30000,0\n29,61000,0\n20,74000,0\n26,15000,0\n41,45000,0\n31,76000,0\n36,50000,0\n40,47000,0\n31,15000,0\n46,59000,0\n29,75000,0\n26,30000,0\n32,135000,1\n32,100000,1\n25,90000,0\n37,33000,0\n35,38000,0\n33,69000,0\n18,86000,0\n22,55000,0\n35,71000,0\n29,148000,1\n29,47000,0\n21,88000,0\n34,115000,0\n26,118000,0\n34,43000,0\n34,72000,0\n23,28000,0\n35,47000,0\n25,22000,0\n24,23000,0\n31,34000,0\n26,16000,0\n31,71000,0\n32,117000,1\n33,43000,0\n33,60000,0\n31,66000,0\n20,82000,0\n33,41000,0\n35,72000,0\n28,32000,0\n24,84000,0\n19,26000,0\n29,43000,0\n19,70000,0\n28,89000,0\n34,43000,0\n30,79000,0\n20,36000,0\n26,80000,0\n35,22000,0\n35,39000,0\n49,74000,0\n39,134000,1\n41,71000,0\n58,101000,1\n47,47000,0\n55,130000,1\n52,114000,0\n40,142000,1\n46,22000,0\n48,96000,1\n52,150000,1\n59,42000,0\n35,58000,0\n47,43000,0\n60,108000,1\n49,65000,0\n40,78000,0\n46,96000,0\n59,143000,1\n41,80000,0\n35,91000,1\n37,144000,1\n60,102000,1\n35,60000,0\n37,53000,0\n36,126000,1\n56,133000,1\n40,72000,0\n42,80000,1\n35,147000,1\n39,42000,0\n40,107000,1\n49,86000,1\n38,112000,0\n46,79000,1\n40,57000,0\n37,80000,0\n46,82000,0\n53,143000,1\n42,149000,1\n38,59000,0\n50,88000,1\n56,104000,1\n41,72000,0\n51,146000,1\n35,50000,0\n57,122000,1\n41,52000,0\n35,97000,1\n44,39000,0\n37,52000,0\n48,134000,1\n37,146000,1\n50,44000,0\n52,90000,1\n41,72000,0\n40,57000,0\n58,95000,1\n45,131000,1\n35,77000,0\n36,144000,1\n55,125000,1\n35,72000,0\n48,90000,1\n42,108000,1\n40,75000,0\n37,74000,0\n47,144000,1\n40,61000,0\n43,133000,0\n59,76000,1\n60,42000,1\n39,106000,1\n57,26000,1\n57,74000,1\n38,71000,0\n49,88000,1\n52,38000,1\n50,36000,1\n59,88000,1\n35,61000,0\n37,70000,1\n52,21000,1\n48,141000,0\n37,93000,1\n37,62000,0\n48,138000,1\n41,79000,0\n37,78000,0\n39,134000,1\n49,89000,1\n55,39000,1\n37,77000,0\n35,5700,0\n36,63000,0\n42,73000,1\n43,112000,1\n45,79000,0\n46,117000,1\n58,38000,1\n48,74000,1\n37,137000,1\n37,79000,1\n40,60000,0\n42,54000,0\n51,134000,0\n47,113000,1\n36,125000,1\n38,50000,0\n42,70000,0\n39,96000,1\n38,50000,0\n49,141000,1\n39,79000,0\n39,75000,1\n54,104000,1\n35,55000,0\n45,32000,1\n36,60000,0\n52,138000,1\n53,82000,1\n41,52000,0\n48,30000,1\n48,131000,1\n41,60000,0\n41,72000,0\n42,75000,0\n36,118000,1\n47,107000,1\n38,51000,0\n48,119000,1\n42,65000,0\n40,65000,0\n57,60000,1\n36,54000,0\n58,144000,1\n35,79000,0\n38,55000,0\n39,122000,1\n53,104000,1\n35,75000,0\n38,65000,0\n47,51000,1\n47,105000,1\n41,63000,0\n53,72000,1\n54,108000,1\n39,77000,0\n38,61000,0\n38,113000,1\n37,75000,0\n42,90000,1\n37,57000,0\n36,99000,1\n60,34000,1\n39,
```

```
54,70000,1\r\n41,72000,0\r\n40,71000,1\r\n42,54000,0\r\n43,129000,1\r\n53,34000,
1\r\n47,50000,1\r\n42,79000,0\r\n42,104000,1\r\n59,29000,1\r\n58,47000,1\r\n46,8
8000,1\r\n38,71000,0\r\n54,26000,1\r\n60,46000,1\r\n60,83000,1\r\n39,73000,0\r\n
59,130000,1\r\n37,80000,0\r\n46,32000,1\r\n46,74000,0\r\n42,53000,0\r\n41,87000,
1\r\n58,23000,1\r\n42,64000,0\r\n48,33000,1\r\n44,139000,1\r\n49,28000,1\r\n57,3
3000,1\r\n56,60000,1\r\n49,39000,1\r\n39,71000,0\r\n47,34000,1\r\n48,35000,1\r\n
48,33000,1\r\n47,23000,1\r\n45,45000,1\r\n60,42000,1\r\n39,59000,0\r\n46,41000,1
\r\n51,23000,1\r\n50,20000,1\r\n36,33000,0\r\n49,36000,1'}
```

Initialize the data frame

```
[4]: df=pd.read_csv("/content/drive/MyDrive/Colab Notebooks/Social_Network_Ads.csv")
```

Perform Data Preprocessing

```
[9]: df.head()
```

```
[9]:   Age  EstimatedSalary  Purchased
0    19             19000           0
1    35             20000           0
2    26             43000           0
3    27             57000           0
4    19             76000           0
```

```
[ ]: df.tail()
```

```
[ ]:   Age  EstimatedSalary  Purchased
395   46             41000           1
396   51             23000           1
397   50             20000           1
398   36             33000           0
399   49             36000           1
```

```
[10]: df
```

```
[10]:   Age  EstimatedSalary  Purchased
0    19             19000           0
1    35             20000           0
2    26             43000           0
3    27             57000           0
4    19             76000           0
..   ...             ...           ...
395   46             41000           1
396   51             23000           1
397   50             20000           1
398   36             33000           0
399   49             36000           1
```

[400 rows x 3 columns]

```
[11]: df.describe()
```

```
[11]:   count    Age  EstimatedSalary  Purchased
count  400.000000      400.000000    400.000000
mean    37.655000     69742.500000     0.357500
std     10.482877     34096.960282     0.479864
min     18.000000     15000.000000     0.000000
25%     29.750000     43000.000000     0.000000
50%     37.000000     70000.000000     0.000000
75%     46.000000     88000.000000     1.000000
max     60.000000    150000.000000     1.000000
```

Identification and Handling of Null Values

```
[6]: df.isnull()
```

```
[6]:   Age  EstimatedSalary  Purchased
0  False             False      False
1  False             False      False
2  False             False      False
3  False             False      False
4  False             False      False
..   ...             ...           ...
395 False             False      False
396 False             False      False
397 False             False      False
398 False             False      False
399 False             False      False
```

[400 rows x 3 columns]

```
[7]: df.isna()
```

```
[7]:   Age  EstimatedSalary  Purchased
0  False             False      False
1  False             False      False
2  False             False      False
3  False             False      False
4  False             False      False
..   ...             ...           ...
395 False             False      False
396 False             False      False
397 False             False      False
398 False             False      False
399 False             False      False
```

```
[400 rows x 3 columns]

[8]: df.isnull().sum()

[8]: Age          0
     EstimatedSalary  0
     Purchased     0
     dtype: int64

[12]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 3 columns):
#   Column          Non-Null Count  Dtype
---  ---
0    Age            400 non-null    int64
1    EstimatedSalary 400 non-null    int64
2    Purchased       400 non-null    int64
dtypes: int64(3)
memory usage: 9.5 KB
```

```
[13]: df.dtypes
```


```
[13]: Age          int64
     EstimatedSalary  int64
     Purchased     int64
     dtype: object
```

Import Seaborn and Matplotlib

```
[14]: import seaborn as sns
     import matplotlib.pyplot as plt
```

```
[17]: df.boxplot()
```

```
[17]: <Axes: >
```



```
[19]: X = df.drop(['Purchased'], axis = 1)
     Y = df['Purchased']
```

Use Logistic regression(Train the Machine) to Create Model:

```
[20]: from sklearn.model_selection import train_test_split
     xtrain, xtest, ytrain, ytest = train_test_split(X, Y, test_size =0.
     ↪2, random_state = 0)
```

```
[21]: ytrain
```

```
[21]: 336    1
     64     0
     55     0
     106    0
     300    1
     ..
     323    1
     192     0
     117     0
     47     0
     172     0
```

Name: Purchased, Length: 320, dtype: int64

```
[22]: xtrain
```

```
[22]:
```

	Age	EstimatedSalary
336	58	144000
64	59	83000
55	24	55000
106	26	35000
300	58	38000
...
323	48	30000
192	29	43000
117	36	52000
47	27	54000
172	26	118000

```
[320 rows x 2 columns]
```

```
[23]: import sklearn
      from sklearn.linear_model import LogisticRegression
```

```
[24]: logreg = LogisticRegression()
```

```
[25]: model=logreg.fit(xtrain, ytrain)
```

```
[26]: LogisticRegression()
```

```
[26]: LogisticRegression()
```

Predict the y_{pred}

```
[27]: y_pred=logreg.predict(xtest)
```

```
[28]: y_pred
```

```
[28]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
            0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
            0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
            0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
[29]: y_pred_xtrain=logreg.predict(xtrain)
```

```
[30]: y_pred_xtrain
```

```
[30]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
            0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

[illegible]

```
[31]: print(xtrain)
      print("-----\n")
      print(xtest)
      print("-----\n")
      print(ytrain)
      print("-----\n")
      print(ytest)
      print("-----\n")
      print(y_pred)
      print("-----\n")
      print(y_pred_xtrain)
```

	Age	EstimatedSalary
336	58	144000
64	59	83000
55	24	55000
106	26	35000
300	58	38000
...
323	48	30000
192	29	43000
117	36	52000
47	27	54000
172	26	118000

```
[320 rows x 2 columns]
```

	Age	EstimatedSalary
132	30	87000
309	38	50000
341	35	75000
196	30	79000

6. Classification Report

```
[40]: print("classification report: ", classification_report(ytest, y_pred))
```

classification report:		precision	recall	f1-score	support
0	0.72	1.00	0.84	58	
1	0.00	0.00	0.00	22	
accuracy			0.73	80	
macro avg		0.36	0.50	0.42	80
weighted avg		0.53	0.72	0.61	80

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```